

Getting connect to the cluster

1. Install Helm 3 and kubectl. Latest versions are fine
 - a. <https://helm.sh/docs/intro/install/>
 - b. <https://kubernetes.io/docs/tasks/tools/install-kubectl-macos/>
2. Login to the purdue VPN to reach geddes
3. Login to geddes and navigate to the geddes cluster with the top left dropdown
4. Click **kubeconfig File** at the top left and copy context to clipboard
5. Create a directory called **.kube** in your homedir (make sure to include the .)
6. Create a file called config in the new **.kube** directory called **config** and copy the kubectl file context here
7. Test to see if kubectl is properly connected to the cluster with the command **kubectl get nodes**

Launching Jupyterhub

1. Create Namespace to deploy Jupyterhub to (remember that the name of your namespace will be part of your domain name in the format **<service-name>.<namespace>.geddes.rcac.purdue.edu** so name appropriately)
 - a. In the Rancher UI Navigate to your ironhacks project using the top left dropdown
 - b. At the top click **namespaces**
 - c. At top left click **add namespace**
 - d. Name and create, leave everything else default
2. Copy context into a file called config.yaml

```
3. singleuser:
4.   image:
5.     name: geddes-registry.rcac.purdue.edu/docker-hub-cache/jupyter/datascience-notebook
6.     tag: latest
7.   networkTools:
8.     image:
9.       name: geddes-registry.rcac.purdue.edu/docker-hub-cache/jupyterhub/k8s-network-tools
10.
11. hub:
12.   image:
13.     name: geddes-registry.rcac.purdue.edu/docker-hub-cache/jupyterhub/k8s-hub
14.
15. proxy:
16.   chp:
17.     image:
18.       name: geddes-registry.rcac.purdue.edu/docker-hub-cache/jupyterhub/configurable-http-proxy
19.       tag: 4.4.0
20.
21. prePuller:
22.   hook:
23.     image:
24.       name: geddes-registry.rcac.purdue.edu/docker-hub-cache/jupyterhub/k8s-image-awaiter
```

3. Add the jupyterhub helm repo with command
 - a. **helm repo add jupyterhub https://jupyterhub.github.io/helm-chart/**
4. Update repo
 - a. **helm repo update**
5. Use the command **helm upgrade --cleanup-on-fail --install <release-name> jupyterhub/jupyterhub --namespace <your-namespace> --values config.yaml** to launch jupyterhub
 - a. **<release-name>** is the release name (this is an arbitrary name to keep track of your helm deployment)
 - b. **<your namespace>** name of the namespace you just created

Domain Name Generation and Service Creation

1. In the Rancher UI create a service
 - a. From the workloads section click on the “service discovery” tab
 - b. Click “add record” at the top
 - c. Give name
 - i. This will act as the service-name portion of the kubernetes dns formatting <service-name>.<namespace>.geddes.rcac.purdue.edu. so in my case my dns name will be postgres.public-test.geddes.rcac.purdue.edu when finished
 - d. Chose the namespace you deployed to.
 - e. Click on “one or more workloads”
 - f. Click on “add target workload”
 - g. Select the **proxy** workload deployed with jupyterhub

- h. Click on “show advanced options” at the bottom right
 - i. It’s a bit hard to see, it’s above the create/cancel button line to the right
- i. Select “layer-4 load balancer” in the “as a” dropdown
 - i. By default this will give your container a campus private address so it can be reached via the DNS name like a normal server deployment but only internally to Purdue. You can leave all the extra configs default that pop up in this section

- ii. If you want to deploy to a public internet address see below steps **l** and **m**
- j. Click “add port”
- k. Fill in “80” under “publish the service port” and “8000” under target port
 - i. This says that connection hitting port 80 (default http port) on your service IP routes to your container port 8000 where the proxy server is listening

The screenshot shows the configuration for a Layer-4 Load Balancer. It includes sections for External Traffic Policy (Local/Cluster), Published IP Addresses (Only ready endpoints/All endpoints), Load Balancer IP (e.g., 78.11.24.19), and Port Mapping. The Port Mapping section shows a port named 'myport' mapping from service port 80 to target port 8000 using the TCP protocol.

- l. (optional public address config) click the **labels & annotations** dropdown and click **add annotation**
- m. (optional public address config) copy **metallb.universe.tf/address-pool: geddes-public-pool** into the key box

The screenshot shows the 'Labels & Annotations' section. It has a table with 'Key' and 'Value' columns. A key 'metallb.universe.tf/address-pool' is entered, and the value 'geddes-public-pool' is entered. There is an 'Add Annotation' button at the bottom.

- n. Click create
2. You should now be able to navigate to your new service URL **<service-name>.<namespace>.geddes.rcac.purdue.edu** and see the jupyterhub login

Updating jupyterhub with more advanced config/auth

1. Update your current deployment by making changes to your current config.yaml and reuse the same command you used initially to launch jupyterhub from section **Launching Jupyterhub** step 5

Documentation

Main jupyterhub docs for customizing deployments:

<https://zero-to-jupyterhub.readthedocs.io/en/latest/jupyterhub/customization.html>

Google auth documentation example:

<https://zero-to-jupyterhub.readthedocs.io/en/latest/administrator/authentication.html?highlight=authentication#google>

How to customize notebook environments:

<https://zero-to-jupyterhub.readthedocs.io/en/latest/jupyterhub/customizing/user-environment.html>